

## Assignment -2

|                     |                     |
|---------------------|---------------------|
| Assignment Date     | 17 September 2022   |
| Student Name        | Sivagama Sundari. V |
| Student Roll Number | 211419104255        |
| Maximum Marks       | 2 Marks             |

### 1. Download the dataset:

```
In [1]: #Downloaded

#importing the libraries

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

import warnings
warnings.filterwarnings('ignore')
```

### 2. Load the dataset.

```
In [2]: #Loading the dataset

d = pd.read_csv(r'Downloads/Churn_Modelling.csv')
```


### 3. Perform Below Visualizations.

#### • Univariate Analysis

```
In [3]: #Boxplot

sns.boxplot(d['CreditScore'])
```


```
Out[3]: <AxesSubplot:xlabel='CreditScore'>
```



In [4]: `#histogram`

```
plt.hist(d['CreditScore'])
```


Out[4]: (array([ 19., 166., 447., 958., 1444., 1866., 1952., 1525., 968.,  
655.]),  
array([350., 400., 450., 500., 550., 600., 650., 700., 750., 800., 850.]),  
<BarContainer object of 10 artists>)



In [5]: `#line plot`

```
plt.plot(d['Age'].head(10))
```


Out[5]: [`<matplotlib.lines.Line2D at 0x1fd895767c0>`]



In [6]: #piechart

```
plt.pie(d['Age'].head(), autopct='%%.2f')
```


Out[6]: (  
 <matplotlib.patches.Wedge at 0x1fd895d4f10>,  
 <matplotlib.patches.Wedge at 0x1fd895e1640>,  
 <matplotlib.patches.Wedge at 0x1fd895e1d60>,  
 <matplotlib.patches.Wedge at 0x1fd895ee4c0>,  
 <matplotlib.patches.Wedge at 0x1fd895eebe0>],  
 [Text(0.8839942345509236, 0.654640506904917, ''),  
 Text(-0.3525952068146547, 1.0419580702366729, ''),  
 Text(-1.09987331875942, -0.01669379169450419, ''),  
 Text(-0.35259525559223215, -1.0419580537304987, ''),  
 Text(0.8739574598774371, -0.6679808068534441, '')],  
 [Text(0.48217867339141285, 0.3570766401299547, '20.29'),  
 Text(-0.19232465826253894, 0.5683407655836397, '19.81'),  
 Text(-0.5999309011415017, -0.009105704560638648, '20.29'),  
 Text(-0.19232468486849025, -0.5683407565802719, '18.84'),  
 Text(0.47670406902405654, -0.3643531673746058, '20.77')])



In [7]: #distplot

```
sns.distplot(d['Age'].head(200))
```

Out[7]: <AxesSubplot:xlabel='Age', ylabel='Density'>




## • Bi - Variate Analysis

In [8]: *#scatter plot*

```
plt.scatter(d['RowNumber'].head(500),d['Age'].head(500))
```

Out[8]: <matplotlib.collections.PathCollection at 0x1fd89b27910>




In [9]: *#bar plot*

```
plt.bar(d['RowNumber'].head(10),d['CreditScore'].head(10))
```

*#labelling of x,y and result*


```
plt.title('Bar plot')  
plt.xlabel('Row Number')  
plt.ylabel('CreditScore')
```

Out[9]: Text(0, 0.5, 'CreditScore')




```
In [10]: #joint plot
sns.jointplot(d['RowNumber'].head(50),d['CreditScore'].head(50))
```

```
Out[10]: <seaborn.axisgrid.JointGrid at 0x1fd89bc0b80>
```




```
In [11]: #bar plot
sns.barplot('RowNumber', 'CreditScore', hue='Geography', data=d.head())
```

```
Out[11]: <AxesSubplot:xlabel='RowNumber', ylabel='CreditScore'>
```



```
In [12]: sns.lineplot(d['RowNumber'].head(),d['CreditScore'].head())
```

```
Out[12]: <AxesSubplot:xlabel='RowNumber', ylabel='CreditScore'>
```




## • Multi - Variate Analysis

```
In [13]: #boxplot
```

```
sns.boxplot(d['RowNumber'].head(10),d['Age'].head(10),d['CreditScore'].head(10))
```


```
Out[13]: <AxesSubplot:xlabel='RowNumber', ylabel='Age'>
```



In [14]: #heat map

```
sns.heatmap(d.head().corr(),annot=True)
```


Out[14]: <AxesSubplot:>




In [15]: #pair plot

```
sns.pairplot(d.head(),hue='CreditScore')
```


Out[15]: <seaborn.axisgrid.PairGrid at 0x2a8e36ae2e0>






```
In [16]: sns.pairplot(d.head())
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x2a8e3329580>
```







## 4. Perform descriptive statistics on the dataset.

In [17]: `#mean`

```
d.mean()
```

```
Out[17]: RowNumber      5.000500e+03
CustomerId    1.569094e+07
CreditScore   6.505288e+02
Age           3.892180e+01
Tenure        5.012800e+00
Balance       7.648589e+04
NumOfProducts 1.530200e+00
HasCrCard     7.055000e-01
IsActiveMember 5.151000e-01
EstimatedSalary 1.000902e+05
Exited       2.037000e-01
dtype: float64
```

In [18]: `#median`

```
d.median()
```

```
Out[18]: RowNumber      5.000500e+03
CustomerId    1.569074e+07
CreditScore   6.520000e+02
Age           3.700000e+01
Tenure        5.000000e+00
Balance       9.719854e+04
NumOfProducts 1.000000e+00
HasCrCard     1.000000e+00
IsActiveMember 1.000000e+00
EstimatedSalary 1.001939e+05
Exited       0.000000e+00
dtype: float64
```

In [19]: `#mode`

```
d.mode()
```

Out[19]:

|      | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age  | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|------|-----------|------------|---------|-------------|-----------|--------|------|--------|---------|---------------|-----------|----------------|-----------------|
| 0    | 1         | 15565701   | Smith   | 850.0       | France    | Male   | 37.0 | 2.0    | 0.0     | 1.0           | 1.0       | 1.0            | 24924           |
| 1    | 2         | 15565706   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 2    | 3         | 15565714   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 3    | 4         | 15565779   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 4    | 5         | 15565796   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| ...  | ...       | ...        | ...     | ...         | ...       | ...    | ...  | ...    | ...     | ...           | ...       | ...            | ...             |
| 9995 | 9996      | 15815628   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 9996 | 9997      | 15815645   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 9997 | 9998      | 15815656   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 9998 | 9999      | 15815660   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |
| 9999 | 10000     | 15815690   | NaN     | NaN         | NaN       | NaN    | NaN  | NaN    | NaN     | NaN           | NaN       | NaN            | N               |

10000 rows × 14 columns

In [20]: `d.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   RowNumber       10000 non-null  int64  
1   CustomerId      10000 non-null  int64  
2   Surname         10000 non-null  object  
3   CreditScore     10000 non-null  int64  
4   Geography       10000 non-null  object  
5   Gender          10000 non-null  object  
6   Age             10000 non-null  int64  
7   Tenure          10000 non-null  int64  
8   Balance         10000 non-null  float64 
9   NumOfProducts  10000 non-null  int64  
10  HasCrCard       10000 non-null  int64  
11  IsActiveMember  10000 non-null  int64  
12  EstimatedSalary 10000 non-null  float64 
13  Exited          10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [21]: `d.shape`

Out[21]: (10000, 14)

In [22]: *#kurtosis*

```
d.kurt()
```

```
Out[22]: RowNumber      -1.200000
CustomerId    -1.196113
CreditScore   -0.425726
Age           1.395347
Tenure        -1.165225
Balance       -1.489412
NumOfProducts 0.582981
HasCrCard     -1.186973
IsActiveMember -1.996747
EstimatedSalary -1.181518
Exited        0.165671
dtype: float64
```

In [23]: *#skewness*

```
d.skew()
```

```
Out[23]: RowNumber      0.000000
CustomerId    0.001149
CreditScore  -0.071607
Age           1.011320
Tenure        0.010991
Balance      -0.141109
NumOfProducts 0.745568
HasCrCard    -0.901812
IsActiveMember -0.060437
EstimatedSalary 0.002085
Exited       1.471611
dtype: float64
```

In [24]: *#standard deviation*

```
d.std()
```

```
Out[24]: RowNumber      2886.895680
CustomerId    71936.186123
CreditScore    96.653299
Age           10.487806
Tenure         2.892174
Balance       62397.405202
NumOfProducts  0.581654
HasCrCard      0.455840
IsActiveMember 0.499797
EstimatedSalary 57510.492818
Exited         0.402769
dtype: float64
```

In [25]: *#variance*

```
d.var()
```

```
Out[25]: RowNumber      8.334167e+06
CustomerId    5.174815e+09
CreditScore   9.341860e+03
Age           1.099941e+02
Tenure        8.364673e+00
Balance       3.893436e+09
NumOfProducts 3.383218e-01
HasCrCard     2.077905e-01
IsActiveMember 2.497970e-01
EstimatedSalary 3.307457e+09
Exited        1.622225e-01
dtype: float64
```

In [26]: `#statistical analysis`

```
d.describe()
```

Out[26]:

|       | RowNumber   | CustomerId   | CreditScore  | Age          | Tenure       | Balance       | NumOfProducts | HasCrCard   | IsActiveMember | EstimatedSalary |   |
|-------|-------------|--------------|--------------|--------------|--------------|---------------|---------------|-------------|----------------|-----------------|---|
| count | 10000.00000 | 1.000000e+04 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000  | 10000.000000  | 10000.00000 | 10000.000000   | 10000.000000    | 1 |
| mean  | 5000.50000  | 1.569094e+07 | 650.528800   | 38.921800    | 5.012800     | 76485.889288  | 1.530200      | 0.70550     | 0.515100       | 100090.239881   |   |
| std   | 2888.89568  | 7.193619e+04 | 96.653299    | 10.487806    | 2.892174     | 62397.405202  | 0.581654      | 0.45584     | 0.499797       | 57510.492818    |   |
| min   | 1.00000     | 1.556570e+07 | 350.000000   | 18.000000    | 0.000000     | 0.000000      | 1.000000      | 0.00000     | 0.000000       | 11.580000       |   |
| 25%   | 2500.75000  | 1.562853e+07 | 584.000000   | 32.000000    | 3.000000     | 0.000000      | 1.000000      | 0.00000     | 0.000000       | 51002.110000    |   |
| 50%   | 5000.50000  | 1.569074e+07 | 652.000000   | 37.000000    | 5.000000     | 97198.540000  | 1.000000      | 1.00000     | 1.000000       | 100193.915000   |   |
| 75%   | 7500.25000  | 1.575323e+07 | 718.000000   | 44.000000    | 7.000000     | 127644.240000 | 2.000000      | 1.00000     | 1.000000       | 149388.247500   |   |
| max   | 10000.00000 | 1.581569e+07 | 850.000000   | 92.000000    | 10.000000    | 250898.090000 | 4.000000      | 1.00000     | 1.000000       | 199992.480000   |   |

In [27]: `#finding unique values for columns`

```
d['Gender'].unique()
```

Out[27]: array(['Female', 'Male'], dtype=object)

In [28]: `d['Geography'].unique()`

Out[28]: array(['France', 'Spain', 'Germany'], dtype=object)

In [29]: `quantile= d['Age'].quantile(q=[0.75, 0.25])`  
`quantile`

Out[29]: 0.75 44.0  
0.25 32.0  
Name: Age, dtype: float64

## 5. Handle the Missing values.

In [30]: `#finding missing values`

```
d.isna()
```

Out[30]:

|      | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age   | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSa |   |
|------|-----------|------------|---------|-------------|-----------|--------|-------|--------|---------|---------------|-----------|----------------|-------------|---|
| 0    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 1    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 2    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 3    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 4    | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| ...  | ...       | ...        | ...     | ...         | ...       | ...    | ...   | ...    | ...     | ...           | ...       | ...            | ...         |   |
| 9995 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 9996 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 9997 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 9998 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |
| 9999 | False     | False      | False   | False       | False     | False  | False | False  | False   | False         | False     | False          | False       | F |

10000 rows × 14 columns

```
In [31]: d.isna().any()
```

```
Out[31]: RowNumber      False
CustomerId    False
Surname        False
CreditScore    False
Geography      False
Gender         False
Age            False
Tenure         False
Balance        False
NumOfProducts  False
HasCrCard      False
IsActiveMember False
EstimatedSalary False
Exited         False
dtype: bool
```

```
In [32]: d.isna().sum()
```

```
Out[32]: RowNumber      0
CustomerId    0
Surname        0
CreditScore    0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

```
In [33]: d.isna().any().sum()
```

```
Out[33]: 0
```


```
In [34]: #no missing values
```

## 6. Find the outliers and replace the outliers

In [35]: `#finding outliers`

```
sns.boxplot(d['Age'])
```

Out[35]: `<AxesSubplot: xlabel='Age'>`



In [36]: `#handling outliers`

```
qnt=d.quantile(q=[0.25,0.75])  
qnt
```

Out[36]:

|      | RowNumber | CustomerId  | CreditScore | Age  | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-----------|-------------|-------------|------|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0.25 | 2500.75   | 15628528.25 | 584.0       | 32.0 | 3.0    | 0.00      | 1.0           | 0.0       | 0.0            | 51002.1100      | 0.0    |
| 0.75 | 7500.25   | 15753233.75 | 718.0       | 44.0 | 7.0    | 127644.24 | 2.0           | 1.0       | 1.0            | 149388.2475     | 0.0    |

In [37]: `iqr=qnt.loc[0.75]-qnt.loc[0.25]`

```
iqr
```

Out[37]:

|                 |             |
|-----------------|-------------|
| RowNumber       | 4999.5000   |
| CustomerId      | 124705.5000 |
| CreditScore     | 134.0000    |
| Age             | 12.0000     |
| Tenure          | 4.0000      |
| Balance         | 127644.2400 |
| NumOfProducts   | 1.0000      |
| HasCrCard       | 1.0000      |
| IsActiveMember  | 1.0000      |
| EstimatedSalary | 98386.1375  |
| Exited          | 0.0000      |

dtype: float64

In [38]: `lower=qnt.loc[0.25]-(1.5*iqr)`  
`lower`

Out[38]:

|                 |               |
|-----------------|---------------|
| RowNumber       | -4.998500e+03 |
| CustomerId      | 1.544147e+07  |
| CreditScore     | 3.830000e+02  |
| Age             | 1.400000e+01  |
| Tenure          | -3.000000e+00 |
| Balance         | -1.914664e+05 |
| NumOfProducts   | -5.000000e-01 |
| HasCrCard       | -1.500000e+00 |
| IsActiveMember  | -1.500000e+00 |
| EstimatedSalary | -9.657710e+04 |
| Exited          | 0.000000e+00  |

dtype: float64


```
In [39]: upper=qnt.loc[0.75]+(1.5*iqr)
upper
```

```
Out[39]: RowNumber      1.499950e+04
CustomerId    1.594029e+07
CreditScore   9.190000e+02
Age           6.200000e+01
Tenure        1.300000e+01
Balance       3.191106e+05
NumOfProducts 3.500000e+00
HasCrCard     2.500000e+00
IsActiveMember 2.500000e+00
EstimatedSalary 2.969675e+05
Exited        0.000000e+00
dtype: float64
```

```
In [40]: #replacing outliers

d['Age']=np.where(d['Age']>57,39,d['Age'])
sns.boxplot(d['Age'])
```

```
Out[40]: <AxesSubplot:xlabel='Age'>
```



## 7. Check for Categorical columns and perform encoding.

```
In [41]: #checking for categorical columns
```

```
d.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   RowNumber             10000 non-null  int64  
 1   CustomerId            10000 non-null  int64  
 2   Surname                10000 non-null  object  
 3   CreditScore            10000 non-null  int64  
 4   Geography              10000 non-null  object  
 5   Gender                 10000 non-null  object  
 6   Age                   10000 non-null  int64  
 7   Tenure                 10000 non-null  int64  
 8   Balance                10000 non-null  float64 
 9   NumOfProducts         10000 non-null  int64  
10   HasCrCard              10000 non-null  int64  
11   IsActiveMember         10000 non-null  int64  
12   EstimatedSalary        10000 non-null  float64 
13   Exited                 10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [42]: d['Geography'].unique()
```

```
Out[42]: array(['France', 'Spain', 'Germany'], dtype=object)
```

```
In [43]: d['Gender'].unique()
```

```
Out[43]: array(['Female', 'Male'], dtype=object)
```

```
In [44]: d['Surname'].unique()
```

```
Out[44]: array(['Hangrave', 'Hill', 'Onio', ..., 'Kashiwagi', 'Aldridge',
              'Burbidge'], dtype=object)
```



In [45]: *#one hot encoding*

```
d['Gender'].replace({'Male':1,'Female':0},inplace=True)
d
```

Out[45]:

|      | RowNumber | CustomerId | Surname   | CreditScore | Geography | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|------|-----------|------------|-----------|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|
| 0    | 1         | 15634602   | Hargrave  | 619         | France    | 0      | 42  | 2      | 0.00      | 1             | 1         | 1              | 1013            |
| 1    | 2         | 15647311   | Hill      | 608         | Spain     | 0      | 41  | 1      | 83807.86  | 1             | 0         | 1              | 1125            |
| 2    | 3         | 15619304   | Onio      | 502         | France    | 0      | 42  | 8      | 159660.80 | 3             | 1         | 0              | 1139            |
| 3    | 4         | 15701354   | Boni      | 699         | France    | 0      | 39  | 1      | 0.00      | 2             | 0         | 0              | 938             |
| 4    | 5         | 15737888   | Mitchell  | 850         | Spain     | 0      | 43  | 2      | 125510.82 | 1             | 1         | 1              | 790             |
| ...  | ...       | ...        | ...       | ...         | ...       | ...    | ... | ...    | ...       | ...           | ...       | ...            | ...             |
| 9995 | 9996      | 15606229   | Obijaku   | 771         | France    | 1      | 39  | 5      | 0.00      | 2             | 1         | 0              | 962             |
| 9996 | 9997      | 15569892   | Johnstone | 516         | France    | 1      | 35  | 10     | 57369.61  | 1             | 1         | 1              | 1016            |
| 9997 | 9998      | 15584532   | Liu       | 709         | France    | 0      | 36  | 7      | 0.00      | 1             | 0         | 1              | 420             |
| 9998 | 9999      | 15682355   | Sabbatini | 772         | Germany   | 1      | 42  | 3      | 75075.31  | 2             | 1         | 0              | 928             |
| 9999 | 10000     | 15628319   | Walker    | 792         | France    | 0      | 28  | 4      | 130142.79 | 1             | 1         | 0              | 381             |

10000 rows × 14 columns

In [46]: *#using dummy variables to encode*

```
d=pd.get_dummies(d,columns=['Geography'])
d
```

Out[46]:

|      | RowNumber | CustomerId | Surname   | CreditScore | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary | Exited |
|------|-----------|------------|-----------|-------------|--------|-----|--------|-----------|---------------|-----------|----------------|-----------------|--------|
| 0    | 1         | 15634602   | Hargrave  | 619         | 0      | 42  | 2      | 0.00      | 1             | 1         | 1              | 101348.88       | 1      |
| 1    | 2         | 15647311   | Hill      | 608         | 0      | 41  | 1      | 83807.86  | 1             | 0         | 1              | 112542.58       | 0      |
| 2    | 3         | 15619304   | Onio      | 502         | 0      | 42  | 8      | 159660.80 | 3             | 1         | 0              | 113931.57       | 1      |
| 3    | 4         | 15701354   | Boni      | 699         | 0      | 39  | 1      | 0.00      | 2             | 0         | 0              | 93826.63        | 0      |
| 4    | 5         | 15737888   | Mitchell  | 850         | 0      | 43  | 2      | 125510.82 | 1             | 1         | 1              | 79084.10        | 0      |
| ...  | ...       | ...        | ...       | ...         | ...    | ... | ...    | ...       | ...           | ...       | ...            | ...             | ...    |
| 9995 | 9996      | 15606229   | Obijaku   | 771         | 1      | 39  | 5      | 0.00      | 2             | 1         | 0              | 96270.64        | 0      |
| 9996 | 9997      | 15569892   | Johnstone | 516         | 1      | 35  | 10     | 57369.61  | 1             | 1         | 1              | 101699.77       | 0      |
| 9997 | 9998      | 15584532   | Liu       | 709         | 0      | 36  | 7      | 0.00      | 1             | 0         | 1              | 42085.58        | 1      |
| 9998 | 9999      | 15682355   | Sabbatini | 772         | 1      | 42  | 3      | 75075.31  | 2             | 1         | 0              | 92888.52        | 1      |
| 9999 | 10000     | 15628319   | Walker    | 792         | 0      | 28  | 4      | 130142.79 | 1             | 1         | 0              | 38190.78        | 0      |

10000 rows × 16 columns

```
In [47]: d=pd.get_dummies(d,columns=['Surname'])
d
```

```
Out[47]:
```

|      | RowNumber | CustomerId | CreditScore | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | ... | Surname_Zinachukwudi | Surna |
|------|-----------|------------|-------------|--------|-----|--------|-----------|---------------|-----------|----------------|-----|----------------------|-------|
| 0    | 1         | 15634602   | 619         | 0      | 42  | 2      | 0.00      | 1             | 1         | 1              | ... | 0                    |       |
| 1    | 2         | 15647311   | 608         | 0      | 41  | 1      | 83807.86  | 1             | 0         | 1              | ... | 0                    |       |
| 2    | 3         | 15619304   | 502         | 0      | 42  | 8      | 159660.80 | 3             | 1         | 0              | ... | 0                    |       |
| 3    | 4         | 15701354   | 699         | 0      | 39  | 1      | 0.00      | 2             | 0         | 0              | ... | 0                    |       |
| 4    | 5         | 15737888   | 850         | 0      | 43  | 2      | 125510.82 | 1             | 1         | 1              | ... | 0                    |       |
| ...  | ...       | ...        | ...         | ...    | ... | ...    | ...       | ...           | ...       | ...            | ... | ...                  |       |
| 9995 | 9996      | 15606229   | 771         | 1      | 39  | 5      | 0.00      | 2             | 1         | 0              | ... | 0                    |       |
| 9996 | 9997      | 15569892   | 516         | 1      | 35  | 10     | 57369.61  | 1             | 1         | 1              | ... | 0                    |       |
| 9997 | 9998      | 15584532   | 709         | 0      | 36  | 7      | 0.00      | 1             | 0         | 1              | ... | 0                    |       |
| 9998 | 9999      | 15682355   | 772         | 1      | 42  | 3      | 75075.31  | 2             | 1         | 0              | ... | 0                    |       |
| 9999 | 10000     | 15628319   | 792         | 0      | 28  | 4      | 130142.79 | 1             | 1         | 0              | ... | 0                    |       |

10000 rows × 2947 columns

## 8. Split the data into dependent and independent variables.

```
In [48]: d.head()
```

```
Out[48]:
```

|   | RowNumber | CustomerId | CreditScore | Gender | Age | Tenure | Balance   | NumOfProducts | HasCrCard | IsActiveMember | ... | Surname_Zinachukwudi | Surname_ |
|---|-----------|------------|-------------|--------|-----|--------|-----------|---------------|-----------|----------------|-----|----------------------|----------|
| 0 | 1         | 15634602   | 619         | 0      | 42  | 2      | 0.00      | 1             | 1         | 1              | ... | 0                    |          |
| 1 | 2         | 15647311   | 608         | 0      | 41  | 1      | 83807.86  | 1             | 0         | 1              | ... | 0                    |          |
| 2 | 3         | 15619304   | 502         | 0      | 42  | 8      | 159660.80 | 3             | 1         | 0              | ... | 0                    |          |
| 3 | 4         | 15701354   | 699         | 0      | 39  | 1      | 0.00      | 2             | 0         | 0              | ... | 0                    |          |
| 4 | 5         | 15737888   | 850         | 0      | 43  | 2      | 125510.82 | 1             | 1         | 1              | ... | 0                    |          |

5 rows × 2947 columns

```
In [49]: x=d.iloc[ : ].values
y=d.iloc[:,2].values
print(x)

[[1.0000000e+00 1.5634602e+07 6.1900000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]
 [2.0000000e+00 1.5647311e+07 6.0800000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]
 [3.0000000e+00 1.5619304e+07 5.0200000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]
 ...
 [9.9980000e+03 1.5584532e+07 7.0900000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]
 [9.9990000e+03 1.5682355e+07 7.7200000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]
 [1.0000000e+04 1.5628319e+07 7.9200000e+02 ... 0.0000000e+00
 0.0000000e+00 0.0000000e+00]]
```

```
In [50]: y
```

```
Out[50]: array([619, 608, 502, ..., 709, 772, 792], dtype=int64)
```

```
In [51]: x=d.drop(columns= ['EstimatedSalary']).values
y=d['EstimatedSalary'].values
x
```

```
Out[51]: array([[1.0000000e+00, 1.5634602e+07, 6.1900000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00],
[2.0000000e+00, 1.5647311e+07, 6.0800000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00],
[3.0000000e+00, 1.5619304e+07, 5.0200000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00],
...,
[9.9980000e+03, 1.5584532e+07, 7.0900000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00],
[9.9990000e+03, 1.5682355e+07, 7.7200000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00],
[1.0000000e+04, 1.5628319e+07, 7.9200000e+02, ..., 0.0000000e+00,
0.0000000e+00, 0.0000000e+00]])
```

```
In [52]: y
```

```
Out[52]: array([101348.88, 112542.58, 113931.57, ..., 42085.58, 92888.52,
38190.78])
```

## 9. Scale the independent variables

```
In [53]: from sklearn.preprocessing import scale #StandardScaler
# Scale - Similar to std
```

```
In [54]: #Scaling the independent variables
```

```
x = scale(x)
x
```

```
Out[54]: array([[ -1.73187761, -0.78321342, -0.32622142, ..., -0.01000005 ,
-0.01414355, -0.01414355],
[ -1.7315312 , -0.60653412, -0.44003595, ..., -0.01000005 ,
-0.01414355, -0.01414355],
[ -1.73118479, -0.99588476, -1.53679418, ..., -0.01000005 ,
-0.01414355, -0.01414355],
...,
[ 1.73118479, -1.47928179, 0.60498839, ..., -0.01000005 ,
-0.01414355, -0.01414355],
[ 1.7315312 , -0.11935577, 1.25683526, ..., -0.01000005 ,
-0.01414355, -0.01414355],
[ 1.73187761, -0.87055909, 1.46377078, ..., -0.01000005 ,
-0.01414355, -0.01414355]])
```

```
In [55]: x.mean()
```

```
Out[55]: 2.348215911527609e-18
```

```
In [56]: x.std()
```

```
Out[56]: 1.000000000000000102
```

## 10. Split the data into training and testing

```
In [57]: from sklearn.model_selection import train_test_split
```

```
In [58]: #splitting data to train and test
```

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2)
print(x_train.shape, x_test.shape)
```

```
(8000, 2946) (2000, 2946)
```